

MGMT 803, MBA@RICE, 2026

Agent security

Kerry Back

J. Howard Creekmore Professor of Finance

This session

- 24. Multi-agent architectures
- 25. Security risks
- 26. Mitigating risks

24 · MULTI-AGENT ARCHITECTURES

Multi-agent teams

“How did our Q3 revenue compare with what analysts expected for the industry?”

Orchestrator

The one you chat with.
Splits the question, then
combines the two answers.

Data agent

Reaches the warehouse.
Returns our Q3 revenue.

Research agent

Reaches the web. Returns
the published industry
estimates.

Calling another agent is a tool use. The orchestrator’s model picks `data_chat` or `web_chat` the way a single agent picks `run_python`.

The tool definition

If an agent is a tool, its definition answers the same three questions as any other tool definition:

Which agents can I call?

The names: `data_chat` and `web_chat`. The model cannot call anything else.

What can I say to them?

The arguments.

`web_chat` takes the user's question only — nothing from the conversation, and so nothing from the data.

What can they give back?

The return. Text, a table, a file — and whatever comes back enters the orchestrator's context.

What the split buys

Separate context windows

Each agent reads only its own material. The orchestrator gets the answer, not the material behind it.

One definition, many callers

The research agent's prompt, tools and examples are written once. Every orchestrator that needs research calls that agent instead of carrying its own copy of the instructions.

A smaller model per job

Retrieval can run on a cheaper model than the one writing the answer.

What else it changes

A shorter tool list

A model choosing among four tools picks better than the same model choosing among forty.

Work at the same time

The data lookup and the web search do not wait for each other.

Tested on its own

You can rewrite the research agent's prompt and test it without retesting the orchestrator.

None of these is a security argument. The security use of the same architecture comes after the risks.

Subagents, or separate agents

A SUBAGENT

The harness starts a second copy of the model with its own context window. It runs on your machine, under your harness, on your keys.

A SEPARATE AGENT

Its own process, its own key, its own MCP connections. The orchestrator reaches it the way it reaches any other tool.

The choice depends on whether the two halves need different credentials.

Breakout exercise

- 1 Build an orchestrator whose only two tools are agents: your FRED agent, and a writing agent that produces a Word document.
- 2 Ask for something that needs both: the last ten years of unemployment and inflation, and a one-page brief on them as a .docx.
- 3 Print which agent was called, with what arguments, and where the file landed.

Work out the design with Claude. Settle what each agent is handed and what it returns before writing anything.

25 · SECURITY RISKS

SENDING DATA TO A PROVIDER

What happens to what you send

It can be trained on

With web search off, ask Claude what Apple's total revenue was in 2020. It answers from the weights, not from a lookup. What goes into a model can come back out of it.

It is retained

Even where nobody trains on it, copies sit on the provider's disks for some window. A breach, a subpoena or an insider reaches them there, under somebody else's security.

Apple's revenue is public. The mechanism is the point: the figure is in the weights, not on a page it fetched.

What regulations require

HIPAA

Satisfactory assurances “must be documented through a written contract” meeting § 164.504(e).

A Business Associates Agreement (BAA) is required. 45 CFR 164.502(e).

SEC REG S-P

Written policies “reasonably designed to require oversight, including through due diligence and monitoring, of service providers.”

17 CFR 248.30.

LETTING IT ACT

Prompt injection

The model cannot tell an instruction from a piece of content. Both arrive as text in the same conversation.

Text placed where the agent will read it — an email, a document, a web page, a spreadsheet cell — is read in the same way as text you typed. An attacker needs no credentials, only a way to get words in front of the model.

OWASP has ranked this first among risks to LLM applications since 2025.

EchoLeak

CVE-2025-32711 · MICROSOFT 365 COPILOT · JUNE 2025

- 1 An attacker mails the victim a message containing instructions invisible to a human reader.
- 2 The victim asks Copilot to summarize their email. Copilot reads the hidden instructions.
- 3 Following them, Copilot places confidential data inside a URL in its response.
- 4 The victim's browser fetches that URL automatically.
- 5 The attacker reads the data out of their own server log.

Zero clicks by the victim beyond the summary request.

Three legs

Private data

Something worth taking:
client records, financials,
an inbox.

Untrusted content

Text an outsider can
control: a web page, an
email, an uploaded file.

An outbound channel

Any way for text to leave:
fetching a URL, sending
mail, writing somewhere
shared.

All three in one agent is the condition. Remove any one and the attack does not work.

26 · MITIGATING RISKS

Training and retention: contractual solutions

	WHAT IT ADDS	ENOUGH FOR
Opt out of training	a setting, revocable, data still retained briefly	personal, non-sensitive work
A signed data processing agreement	liability, breach notice, audit rights, a subprocessor list, zero retention on request	client data, regulated data
That, plus a business associate agreement	the same, covering protected health information	patient records
An enterprise agreement	service levels, support, single sign-on, indemnification, capacity	scale and operations

Data exfiltration: removing a leg

LEG	HOW YOU REMOVE IT	WHAT IT COSTS
Private data	the agent sees only what you hand it — Cowork cannot reach outside its project folder	it cannot answer from data it never sees
Untrusted content	no web fetch, no reading files people upload, no general code execution	a narrow agent, and no <code>run_python</code>
A way out	an egress allowlist, and no tool that sends	results leave only when a person moves them

Any of the three can go. Each takes capability with it, and which one you can afford to lose is the decision.

A WORKED EXAMPLE

Three agents

Orchestrator

The one you chat with. Its only two tools are `web_chat` and `data_chat`. It holds no data and reaches nothing itself.


Webbot

WebSearch and WebFetch, running at Anthropic. It reads untrusted content and holds no private data.

Databot

`run_python` and the MCP data connectors, behind an egress proxy. It holds the private data and cannot reach the web.

The architecture



The web channel carries only the user's words; private data can flow into the data agent, never out toward the web.

What each agent may be told

web_chat IS LOCKED

Its tool definition says it sends the user's verbatim prompt and nothing else. Nothing from the conversation, and so nothing from the data, can reach the agent that can reach the web.

The argument is about the arguments: what a tool accepts decides what can leave through it.

data_chat IS OPEN


The model composes this prompt from the conversation. That is safe only because databot has no route out.

Four containers, four networks

	NETWORKS	WHAT THAT ALLOWS
Orchestrator	corporate, internalA, internalB	serves the chat UI, talks to both agents
Webbot	internalB, public	talks to the orchestrator, calls Anthropic
Databot	internalA	talks to the orchestrator and the proxy, nothing else
Squid proxy	internalA, public	the only route out of the data half

Webbot shares a network with the orchestrator, so the orchestrator’s endpoints must authenticate at the application, not only at an external gateway.

Inside the data half



The databot has no route out except the proxy; nginx lets the user in, and the proxy reaches only the Claude API and the two connectors.

Mechanical guards

Sanitize the exit

Even with the legs split, a delivered file could carry a planted remote image — the EchoLeak pattern. Every file the orchestrator hands back is checked first for anything that auto-fetches an external URL.

Safeguard memory

Memory persists into every later conversation, so a planted instruction would reload forever. Every saved fact must be grounded in the user's own words, not in tool results or fetched content.