

MGMT 803: Gen AI for Business

Syllabus · Fall 2026



Instructor

[Kerry Back](#)

J. Howard Creekmore Professor of Finance and Professor of Economics

kerryback@gmail.com

Meeting Times and Locations

Day	Times	Room
Thursday, September 24	8:00 AM – 6:00 PM	McNair Hall 316
Friday, September 25	8:00 AM – 12:30 PM	McNair Hall 316
Saturday, September 26	8:00 AM – 6:00 PM	McNair Hall 317
Sunday, September 27	8:00 AM – 12:30 PM	McNair Hall 317

Course Description

Generative AI moved into ordinary business work faster than any recent technology. McKinsey's global survey found the share of organizations regularly using generative AI in at least one business function rose from roughly a third in 2023 to about two-thirds a year later, and higher still since. The Census Bureau's Business Trends and Outlook Survey, which asks the narrower question of whether a firm used AI to produce a good or service in the past two weeks, shows the rate roughly doubling over the same period, with adoption concentrated in information and professional services. The gap between those two numbers is itself informative: many people are using these tools, and far fewer organizations have built anything durable around them.

Business uses fall into a few broad categories. There is individual productivity — the drafting, summarizing, analysis, and coding that one person does faster with a model beside them. There is internal deployment — putting a model to work on the firm's own documents and systems, in workflows that several people depend on. There are customer-facing applications, where the model

talks to people outside the firm. And there is the reshaping of the product itself, where AI is what is being sold rather than what makes it. The categories differ less in the technology than in what goes wrong: a bad answer you catch yourself is a nuisance, a bad answer inside a workflow is a process failure, and a bad answer in front of a customer is a liability.

This course covers the first two. We start with individual productivity, because the judgment you need there — knowing what the model is doing, what it is good at, how it fails, and how to check it — is the same judgment every larger deployment rests on. We then move to work that outlives the session you did it in: writing a repeated task down so others can run it, connecting a model to data it was never trained on, generating the document or application someone else opens, and assembling an agent with a deliberately chosen list of things it can do. The course closes on the questions a firm has to answer before switching any of this on — where the data goes, what an outside document can talk the system into doing, and what it costs to run.

The course is hands-on throughout. You work in your own container on the course lab server, and you show what you built to the room.

Learning Objectives

By the end of the course, you should be able to:

1. Work fluently with generative AI in your own tasks — choosing the right tool for the job, directing it well, and verifying what it gives back, with a working understanding of how these models produce their answers and where that mechanism fails.
2. Turn ad hoc use into something durable — a documented procedure, a connection to your own data, a deliverable someone else opens, or an agent with a defined job — that a colleague can run without you.
3. Evaluate a proposed deployment on the terms a firm actually decides on: which model and where it runs, what data it touches, what it is permitted to do, and what it costs relative to what it saves.

Topic Schedule

	When	Topics
1	Thu, Sep 24 8:00 AM – 12:30 PM	1 How a language model works 2 Chat, CLI, and sandbox 3 Prompting tips
2	Thu, Sep 24 1:30 – 6:00 PM	4 Persistent instructions and memory 5 Connectors and plugins 6 Agent harnesses
3	Fri, Sep 25 8:00 AM – 12:30 PM	7 Working with data
4	Sat, Sep 26 8:00 AM – 12:30 PM	8 Creating docs 9 Creating apps
5	Sat, Sep 26 1:30 – 6:00 PM	10 Retrieval augmented generation 11 Open source models

	When	Topics
6	Sun, Sep 27	12 Creating agents
	8:00 AM – 12:30 PM	13 Managing security risks

Grading

Grades will be based on in-class exercises, screen shares, and class participation.

Honor Code

The Rice University honor code applies to all work in this course.

Disability Accommodations

Any student with a documented disability requiring accommodations in this course is encouraged to contact me outside of class. All discussions will remain confidential. Any adjustments or accommodations regarding assignments or the final exam must be made in advance. Students with disabilities should also contact Disability Support Services in the Allen Center.